

NOETHER Paper Creation Process Record / 论文创作过程记录

Manuscript / 论文: NOETHER —A Constructive Framework for Metamorphic Pattern Discovery from Operator Algebras **Target venue / 目标刊物:** ACM Transactions on Software Engineering and Methodology (TOSEM) **Final state / 最终状态:** Accept (camera-ready) —42 pages, 853 KB PDF, 54 cited references, 6 supplementary directories **Document version:** 2026-05-03 **Pipeline orchestrator:** academic-pipeline v3.2.2 (Stages 1 → 6)

1. One-page summary / 一页摘要

EN. This document records the 4-round collaboration that produced the NOETHER manuscript. The work began as a theoretical proposal and evolved through four reviewer rounds—Major → Major → Minor → Accept—into a TOSEM-ready submission with 54 citations, 6 supplementary directories, an executed empirical pilot ($n = 5$), and a 13-row revision-provenance table tracing every reviewer concern to its resolution. The framework’s central methodological claim (“use versioned hypotheses + mechanical downstream + transparent revision”) shaped the manuscript’s own production: each round closed previously open items rather than opening new ones, and the manuscript’s authors did not over-claim, did not silently drop concerns, and did not skip integrity checks.

中文. 本文档记录了产生 NOETHER 论文的 4 轮协作过程。工作始于一项理论命题, 经过四轮审稿 (Major → Major → Minor → Accept) 演变为可投稿 TOSEM 的稿件, 包含 54 条引用、6 个补充材料目录、一项已执行的实证 pilot ($n = 5$)、以及一张 13 行修订溯源表, 把每条审稿关切都映射到对应的解决方案。框架的核心方法论主张 (“用版本化假设 + 机械下游 + 透明修订”) 也塑造了稿件自身的生产过程: 每一轮都关闭了上一轮的开放项, 作者没有过度主张、没有悄悄略过关切、没有跳过完整性检查。

2. Revision journey / 修订旅程

Round 1 —TOSEM Major Revision (initial review)

EN. Three core concerns: (1) Theorem 1’s “constructive completeness” is near-circular; (2) §6 equivariant ML transfer is too thin (75% non-empty patterns covered, $|\text{denominator}|=4$, no comparative validation); (3) seven-block decomposition has a circularity in the strong reading of “prediction”— T^* and $\mathcal{T}^*$ blocks were induced from reactor physics, then “predict” reactor-physics MRs.

Author response strategy. Did *not* attempt to prove Theorem 1’ (which would require resolving the open conjecture). Instead, took paths (b) renaming + wording calibration, and (c) explicit out-of-scope MR cataloguing. Added Hypothesis 1 v1.0 framing, Boundary-of-contribution boxed statements, and the §6.6 small-scale comparative case study with real EGNN training and a pre-registered hypothesis check.

Output deliverables. Round-1 response (Response_to_Reviewers.md, 393 lines), supplementary S1–S4, paper grew 27 → ~30 pages.

中文. 三项核心关切: (1) Theorem 1 的“构造性完备性”接近循环; (2) §6 等变 ML 跨域过弱 (75% non-empty patterns covered, 分母仅 4, 无对比验证); (3) 七块分解中“prediction”的强读法存在循环—— T^* 与 $\mathcal{T}^*$ 块是从反应堆物理归纳来的, 再用它“预测”反应堆物理 MR。作者策略: 不尝试证明 Theorem 1’ (等同解决开放猜想); 采取 (b) 重命名 + 措辞校准, (c) 明确 out-of-scope MR 编目。加入 Hypothesis 1 v1.0 框架、Boundary-of-contribution 框注、以及 §6.6 真实 EGNN 训练 + 预登记假设检查的小规模对比 case study。

Round 2 —TOSEM Major Revision (escalated)

EN. New concerns: (1) empirical evaluation insufficient; (2) constructive claim self-undermined by author’s own caveats; (3) transferability under-evidenced because both instantiated domains share Lie-group/self-adjoint/time-reversal mathematical core. Reviewer suggested 12 specific improvements, prominently a third domain instantiation and shared-benchmark comparison.

Strategy. Implemented two-layer claim separation (downstream mechanical vs upstream empirical), explicit Boundary boxes at §1, §3.9, §8, Hypothesis 1 v1.0 versioning. Added §6.7 relational query optimisers as the third domain

—a *productive falsification* that motivates v1.1 with an additional block $\mathcal{B}_{\text{rel}}^*$. Added Noether-style derivations in §5.4 (Boltzmann adjoint) and §6.4 (SO(3) rotation) to make the Noether title link substantive rather than rhetorical. Added 12 new references identified through paper-search-mcp (Murphy 2008, Liu 2014, Nolasco 2024 MemoRIA, Humbatova 2021 DeepCrime, etc.).

Output deliverables. Round-3 response (Response_to_Reviewers_Round3.md, 319 lines), §6.7 (~2 pages new), Definition 14 ($\mathcal{B}_{\text{rel}}^*$), Hypothesis 1 v1.1, Theorem 1/2 v1.1 status remarks. Paper grew 30 → 35 pages.

中文. 新关切:(1) 经验评估不足;(2) 构造性主张被作者自身论述削弱;(3) Transferability 证据不足——两个实例化共用 Lie 群/自伴/时间反演数学骨架。审稿人提了 12 项具体建议, 核心是第三个领域 + 共享 benchmark 比较。策略: 实施双层主张分离 (下游 mechanical vs 上游 empirical)、§1/§3.9/§8 三处 Boundary 框、Hypothesis 1 v1.0 版本化标签。加入 §6.7 关系数据库查询优化器作为第三领域——一次 *productive falsification*, 推动 v1.1 加入第八块 $\mathcal{B}_{\text{rel}}^*$ 。加入 Noether 风格推导 (§5.4 Boltzmann 伴随、§6.4 SO(3) 旋转), 让题目里的“NOETHER”承担方法论而非修饰角色。新增 12 条文献 (经 paper-search-mcp 检索)。

Round 3 —TOSEM Minor Revision

EN. Reviewer upgraded recommendation from Major to Minor. Five priority items: (1) execute at least one comparative protocol; (2) upgrade §6.7 to v1.1 minimal instantiation; (3) reduce dependence on anonymous [1]/[2]; (4) Abstract / §1 case-study claim caveats; (5) §7.1 R1/R2/R3 mapping. Five detail items.

Strategy. Critical: executed an actual $n = 5$ DeepCrime-style real-fault pilot on the trained EGNN checkpoint (5 mutation operators systematically derived from the DeepCrime taxonomy: LR / ACH / LRM / BR / WCI). Pilot results: Set N detects 2/5 (cat-v-01, cat-v-03 via ρ_{train}), Set L 0/5, Set B 0/5. Wilson 95% CIs reported, Fisher-exact $p = 1.00$ at this sample size. **Did not over-interpret this result.** The pilot exposed a new v1.0/v1.1-uncovered class ($\mathcal{B}_{\text{wd}}^*$) that the manuscript catalogues honestly as a v1.2 placeholder.

Also: §6.7 fully formalised ($\mathcal{B}_{\text{rel}}^*$ Definition 14, Table 5 row, Definition 13 v1.1 ordering, Theorem 1/2 v1.1 status remarks). Appendix B added with per-MR independent provenance for all 12 representative MRs of Table 3 (Bell-Glasstone, Lewis-Miller, IAEA, ANS standards) —making §5.3 evaluable independently of [1]/[2]. §7.1 13-row revision-provenance table.

Output deliverables. Round-4 response (Response_to_Reviewers_Round4.md, 239 lines), §6.6.1 + Table 6 (DeepCrime pilot), §7.4 engineering guidance (K-sweep audit, tolerance selection), Appendix B (3 pages, 13 MRs). Paper grew 35 → 40 pages.

中文. 审稿人从 Major 上调为 Minor. 五项优先级:(1) 执行至少一项对比协议;(2) 升级 §6.7 到 v1.1 最小实例化;(3) 降低 [1]/[2] 依赖;(4) Abstract / §1 caveat;(5) §7.1 R1/R2/R3 映射。关键执行: **真跑** $n = 5$ DeepCrime-style 真实故障 pilot 在已训练 EGNN checkpoint 上 (5 个 DeepCrime 算子: LR / ACH / LRM / BR / WCI)。结果:Set N 检出 2/5,L=0/5,B=0/5。Wilson 95% CI 报告,Fisher 精确 $p = 1.00$ 。不过度解读。Pilot 暴露 v1.0/v1.1 未覆盖的新类 ($\mathcal{B}_{\text{wd}}^*$), 诚实地标注为 v1.2 占位。§6.7 完整 v1.1 形式化、附录 B(12 条 MR 独立溯源)、§7.1 13 行修订溯源表。

Round 4 —TOSEM Accept (with 8 camera-ready polish items)

EN. Reviewer recommends Accept. Eight polish items: (1) $\mathcal{B}_{\text{wd}}^*$ formalisation depth; (2) Abstract acknowledge §6.6.1 is *second* adversarial test; (3) K-sweep $\pm 5\%$ threshold + $\tau \approx 100 \epsilon_{\text{fp}}$ justifications; (4) Bibliography count consistency (Table 6 says 55, actually 43 cited); (5) §6.6.1 cat-v-01 detection mechanism explanation; (6) Abstract “matching prediction”precise wording; (7) Noether 1918 Tavel translation DOI; (8) Appendix B (xii.a)/(xii.b) numbering consistency.

Strategy. All 8 items addressed: $\mathcal{B}_{\text{wd}}^*$ given a v1.2 placeholder Definition with explicit out-of-scope framing; Abstract acknowledges both adversarial tests (§6.7 deliberate, §6.6.1 unsolicited); K-sweep cited Lewis-Miller §6.2, τ cited Higham 2002 ($n \cdot \gamma_n$ bound for $n \approx 10^3$ ops); 12 previously uncited references inserted in §2.1, §2.3, §2.4, §6.1; Murphy duplicate removed; cat-v-01 detection mechanism explained (head magnitude collapse → softmax softening → boundary-input argmax drift → ρ_{train} inference-stability fail); Abstract sentence rewritten to match §1 contribution C4 wording; Noether DOI added; Appendix B Group 6 renumbered to (xii)/(xiii).

Output deliverables. Round-5 response (Response_to_Reviewers_Round5.md, 139 lines), Higham 2002 reference added, paper grew 40 → 42 pages, all 54 bib entries cited.

中文. 评审推荐 Accept. 8 项 polish:(1) $\mathcal{B}_{wd}^*$ 形式化深度;(2) Abstract 承认 §6.6.1 是第二处对抗性测试;(3) K-sweep $\pm 5\%$ 阈值 + $\tau \approx 100 \epsilon_{fp}$ 加 justification;(4) bib 数字一致性 (Table 6 说 55 实际 43);(5) §6.6.1 cat-v-01 检测机制解释;(6) Abstract “matching prediction”精确措辞;(7) Noether 1918 Tavel DOI;(8) 附录 B (xii.a)/(xii.b) 编号一致性。8 项全部落地; 详见 Round-5 response。论文 40 $\rightarrow$ 42 页, 54/54 全部被引用。

3. Collaboration quality evaluation / 协作质量评估

Honest scoring across 6 dimensions on a 0–100 rubric. No inflation. Evidence cited inline.

Dimension / 维度	Score	Justification / 评分依据
Research depth / 研究深度	88 / 100	Paper makes a non-trivial theoretical contribution (Theorem 1 closure under Translate, Theorem 2 polynomial-time decidability, CONSTRUCT-MP algorithm). The seven-block decomposition is empirical curation but the framework is honest about this. The downstream mechanical layer is a real CS contribution. Deduction: Theorem 1’ (absolute completeness) remains open.
Writing quality / 写作质量	84 / 100	Three boundary-of-contribution boxes (§1, §3.9, §8) are aligned and consistent. Hypothesis 1 versioning is rigorous. Two-layer claim structure is well-executed. Deduction: 42-page length is at the upper edge of acmsmall envelope; some sections could be more concise.
Methodology rigor / 方法论严谨度	86 / 100	Pre-registered H1/H2 hypotheses, Wilson 95% CIs, Fisher exact tests, McNemar tests for paired comparisons, R1/R2/R3 revision-provenance table. The DeepCrime-style pilot ($n = 5$) is honest about its statistical power. Deduction: full GenMorph 23-Java + DeepCrime real-fault and IMDb / Segura QBS-MR comparisons are still protocols, not results.

Dimension / 维度	Score	Justification / 评分依据
Citation integrity / 引用完整性	94 / 100	All 54 bib entries cited (verified via Python diff at Round-5). Round-3's 12 uncited entries were a real bug and Round-4 caught it; Round-5 fixed it. Bell-Glasstone, Lewis-Miller, IAEA, ANS standards properly anchor §5.3 claims. Deduction: 2 anonymised companion papers; de-anonymisation deferred to acceptance.
Responsiveness to feedback / 响应反馈	92 / 100	All 4 review rounds addressed every reviewer item. R3 minor revision delivered $n = 5$ pilot when R2 only required protocol—reviewer explicitly noted this exceeded expectation. R4 polish items 100% addressed. Deduction: in early rounds the response document had naming inconsistency.
Self-reflection quality / 自省质量	89 / 100	The framework methodologically argues for “use versioned hypotheses + transparent revision”; the manuscript's production process embodies this. R3's pilot exposed $\mathcal{B}_{\text{wd}}^*$ as a v1.0/v1.1-uncovered class; R5 acknowledged this as a v1.2 placeholder. R5 took the <i>harder</i> option for $\mathcal{B}_{\text{wd}}^*$ rather than the easier “out of scope” framing. Deduction: in earlier rounds, the manuscript over-emphasised cross-domain transferability before the third domain instantiation existed.

Composite score / 综合得分: $(88 + 84 + 86 + 94 + 92 + 89)/6 = 88.83$ —qualitatively at the **high end of TOSEM acceptance threshold**.

4. AI self-reflection report (7-mode failure-mode audit) / AI 自省报告

Following the academic-pipeline v3.2 mandatory checklist (run at Stage 4.5 final integrity).

Failure mode / 失效模式	Status	Evidence / 依据
1. Citation hallucination	PASS	All 54 bib entries verified to exist. Three were verified via DOI; rest via paper-search-mcp Consensus database lookup. Anonymous [1]/[2] are author working papers, properly marked @unpublished.

Failure mode / 失效模式	Status	Evidence / 依据
2. Implementation bugs	PASS	The §6.6.1 pilot infrastructure was executed end-to-end against the trained EGNN checkpoint with deterministic seeds. 75-row results CSV produced. Fingerprint-based mutation verification confirms each mutation actually changes the model.
3. Hallucinated results	PASS	Pilot detection counts (N: 2/5, L: 0/5, B: 0/5) come from real torch forward passes, not from fabricated tables. The detection mechanism explanation (R5.5) was added <i>after</i> observing the actual cat-v-01 detection in the result CSV — post-hoc but mechanistically grounded.
4. Shortcut reliance	PASS	The manuscript explicitly identifies and reports its shortcuts: §6.6 case study mutations are hand-constructed (acknowledged in construct-validity caveat); EGNN is a “minimal stand-in” for full SE(3)-Transformer (acknowledged); cat-v-02/04/05 not detected (acknowledged as $\mathcal{B}_{\text{wd}}^*$). No silent shortcuts.
5. Bug-as-insight	PASS	The “framework boundary” framing (no MR set detects category-(i) wrong-sign mutations) is <i>not</i> a bug-as-insight: the underlying observation (label-consistency block missing from v1.0) is an honest framework limitation that motivates v1.2, not a re-spinning of a code defect.
6. Methodology fabrication	PASS	Theorem 1’s proof catalogues every block-block interaction; Theorem 2’s bound is derived from per-block t_i values in Table 1. CONSTRUCT-MP algorithm is reference-implementation-backed (supplementary S1, 13/13 unit tests).
7. Pipeline-level frame-lock	PASS	The manuscript was revised through 4 review rounds without frame-lock: each round’s evaluator was distinct, and the revisions adapted to each round’s specific framing. R3’s H1 demotion to “consistency check”+ introduction of falsifiable H1* is the canonical anti-frame-lock move.

Override count: 0. No reviewer concern was overridden without resolution.

5. Artifact inventory / 产出清单

5.1 Primary deliverables / 主要产出

Artifact / 产出物	Path	Size
Final manuscript PDF	NOETHER_paper.pdf	42 pages, 853 KB
LaTeX source	NOETHER_paper.tex	1263 lines
Bibliography	NOETHER_paper.bib	54 entries (all cited)

5.2 Response documents / 应答文档

Round	File	Lines	Purpose
R1	Response_to_Reviewers1.md	393	First major-revision response
R2	Response_to_Reviewers1_Round3.md	319	Second major-revision response
R3	Response_to_Reviewers3_Round4.md	239	Minor-revision response with $n = 5$ pilot
R4	Response_to_Reviewers3_Round5.md	139	Camera-ready polish response (8 items)

5.3 Supplementary archive / 补充材料

ID	Path	Contents
S1	supplementary/S1_construct_references	CONSTRUCT-MP reference implementation (Python, 13/13 tests)
S2	supplementary/S2_pwr_corpus	84-MR PWR corpus + independent_citation_provenance.md
S3	supplementary/S3_case_study	20-mutation case study + 5-mutation DeepCrime-style pilot + EGNN checkpoint
S4	supplementary/S4_reproduction	scd3, conda env, dataset versions, compute environment
S5	supplementary/S5_real_bug	10 e3nn / PyG real-bug evaluation (protocol-pending)
S6	supplementary/S6_query_opt	NOETHER v1.1 vs Segura QBS-MR (protocol-pending)
Anchor	supplementary_README.md	SHA-256: dc54d8288205c98e1edd2a96e724cdc9261

5.4 Process record / 过程记录

File	Purpose
NOETHER_process_summary.md (this file)	Bilingual 4-round collaboration record + scoring + AI self-reflection
NOETHER_final_integrity_log.md	Stage 4.5 final integrity log

6. Methodological lesson / 方法论启示

EN. The NOETHER manuscript advocates a methodological move: replace inductive grounding with algebraic grounding in the downstream layer, while keeping the upstream layer empirical and honest. The 4-round revision process *itself* embodied this move: each round, claims were either *established* (e.g., Theorem 1, Theorem 2, $\mathcal{B}_{\text{rel}}^*$ Definition 14) or *explicitly scoped open* (Theorem 1', Hypothesis 1's sufficiency, $\mathcal{B}_{\text{wd}}^*$ as v1.2 placeholder), but never silently dropped. The reviewer's R4 verdict —“本稿是方法论与执行高度同构的稿件”—is the closest external observation that this isomorphism was achieved.

中文. NOETHER 论文倡导一种方法论: 用代数性下游 + 经验性上游 + 诚实标注开放项, 替换纯经验归纳。4 轮修订过程本身也体现了这一方法论: 每一轮的主张要么 被建立 (如 Theorem 1、Theorem 2、 $\mathcal{B}_{\text{rel}}^*$ Definition 14), 要么 被明确限定为开放 (Theorem 1'、Hypothesis 1 充分性、 $\mathcal{B}_{\text{wd}}^*$ 作为 v1.2 占位), 从来没有被悄悄略过。审稿人在 R4 判词中所说“本稿是方法论与执行高度同构的稿件”是这一同构性获得外部认可的最接近表述。

Document version: 2026-05-03. Generated by academic-pipeline Stage 6 (PROCESS SUMMARY) following Stage 4.5 (FINAL INTEGRITY) PASS.